



Introduction to MPI and Domain decomposition

Ji-Hoon Kang, jhkang@kisti.re.kr

Korea Institute of Science and technology (KISTI)











Contents

- ▶ Introduction to MPI
- ▶ Point-to-point communication
- ▶ **Collective communication**
- ▶ **Loop parallelization**
- ▶ **Domain decomposition**

Review with examples



Collective communication – schematics





Mores on collective communication

`comm.Gatherv(sbuf, rbuf=(rbuf, scount), root)`

`comm.Scatterv(sbuf=(sbuf,scount), rbuf, root)`

`comm.Allgatherv(sbuf, rbuf=(rbuf, scount))`

`comm.Alltoallv(sbuf=(sbuf,scount), rbuf=(rbuf, scount))`

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD


size = comm.Get_size()
rank = comm.Get_rank()

isend = np.array([1, 2, 2, 3, 3, 3])
irecv = np.zeros(3 * (rank + 1), dtype = int)
iscnt = np.array([1, 2, 3])
ircnt = np.full(3, rank + 1, dtype = int)
isend += size * rank


comm.Alltoallv((isend, iscnt), (irecv, ircnt))
print('Rank({0}) : isend = {1}'.format(rank, isend))
print('Rank({0}) : irecv = {1}'.format(rank, irecv))
```



Gatherv (lab11)



`comm.Gather(sbuf, rbuf, root=0)`



```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD

size = comm.Get_size()
rank = comm.Get_rank()

ircnt = np.array([1, 2, 3], dtype = int)
irecv = np.zeros(6, dtype = int)
isend = np.zeros(rank + 1, dtype = int)

for i in range(rank + 1) :
    isend[i] = rank + 1


comm.Gatherv(isend, (irecv, ircnt), 0)
print('rank = {0}, isend = {1}'.format(rank, isend))

if rank == 0 :
    print('rank = {0}, irecv = {1}'.format(rank, irecv))
```


```
$ mpirun -np 3 python3 lab11_gatherv.py
rank = 1, isend = [2 2]
rank = 2, isend = [3 3 3]
rank = 0, isend = [1]
rank = 0, irecv = [1 2 2 3 3 3]
```



Scatterv (lab16)



`comm.Scatter(sbuf, rbuf, root=0)`



```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD

size = comm.Get_size()
rank = comm.Get_rank()

irecv = np.zeros(rank + 1, dtype = int)
iscnt = np.array([1, 2, 3])


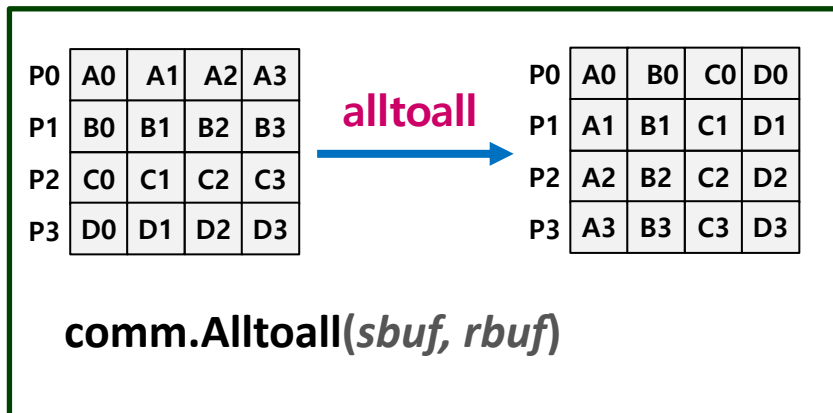
if rank == 0 :
    isend = np.array([1, 2, 2, 3, 3, 3], dtype = int)
else :
    isend = np.zeros(6, dtype = int)

comm.Scatterv([isend, iscnt], irecv, 0)
if rank == 0 :
    print('Before : rank = {0}, irecv = {1}'.format(rank, isend))
    print('After 1 : rank = {0}, irecv = {1}'.format(rank, irecv))
```

```
$ mpirun -np 3 python3 lab16_scatterv.py
After 1 : rank = 1, irecv = [2 2]
Before : rank = 0, irecv = [1 2 2 3 3 3]
After 1 : rank = 0, irecv = [1]
After 1 : rank = 2, irecv = [3 3 3]
```



Alltoallv (lab19)



```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD

size = comm.Get_size()
rank = comm.Get_rank()

isend = np.array([1, 2, 2, 3, 3, 3])
irecv = np.zeros(3 * (rank + 1), dtype = int)
iscnt = np.array([1, 2, 3])
ircnt = np.full(3, rank + 1, dtype = int)
isend += size * rank

comm.Alltoallv((isend, iscnt), (irecv, ircnt))
print('Rank({0}) : isend = {1}'.format(rank, isend))
print('Rank({0}) : irecv = {1}'.format(rank, irecv))
```


```
mpirun -np 3 python3 lab19_alltoallv.py
Rank(0) : isend = [1 2 2 3 3 3]
Rank(0) : irecv = [1 4 7]
Rank(1) : isend = [4 5 5 6 6 6]
Rank(1) : irecv = [2 2 5 5 8 8]
Rank(2) : isend = [7 8 8 9 9 9]
Rank(2) : irecv = [3 3 3 6 6 6 9 9 9]
```


Domain decomposition



Domain decomposition


► 1D decomposition







Domain decomposition

► 2D decomposition





Decomposition type

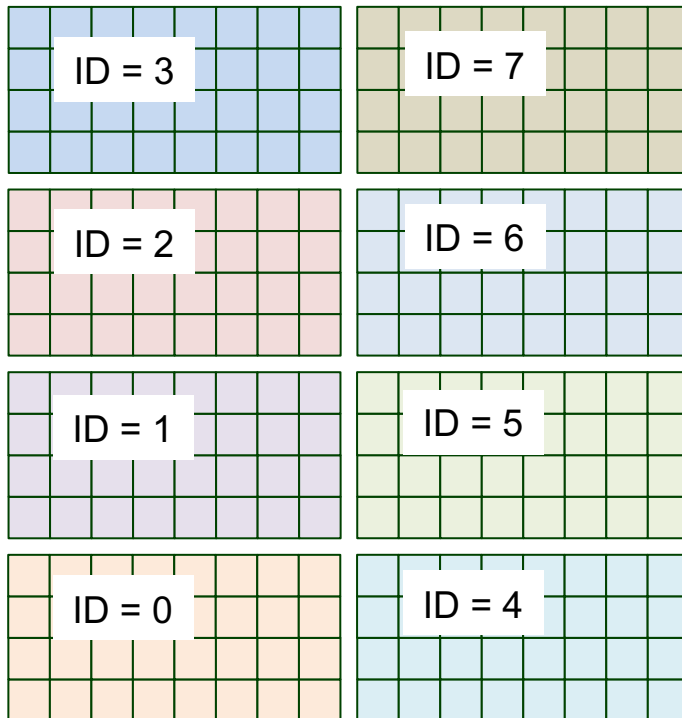


	1-D decomposition	2-D decomposition
Communication pattern	One boundary cells	Two boundary cells
Implementation	Relatively simple	Relatively complicated
Available MPI processes	N_x (or N_y)	$N_x \times N_y$
Communication amount	$2 N_y$ (or $2 N_x$)	$\sim 2(N_x + N_y)/\text{sqrt}(p)$



Mapping domains to processes (I)

- We can design our own methods




For the process of ID=5,
w_ID=1
e_ID=-1
s_ID=4
n_ID=6

OR

For the process of ID=5,
neighbor(4) = (1, -1, 4, 6)

Mapping domains to processes (II)


► Arbitrary numbering



ID	Xp	Xm	Yp	Ym
0	-1	5	7	4
1	4	-1	5	3
2	7	-1	-1	5
3	6	-1	1	-1
4	-1	1	0	6
5	0	-1	2	1
6	-1	3	4	-1
7	-1	2	-1	0

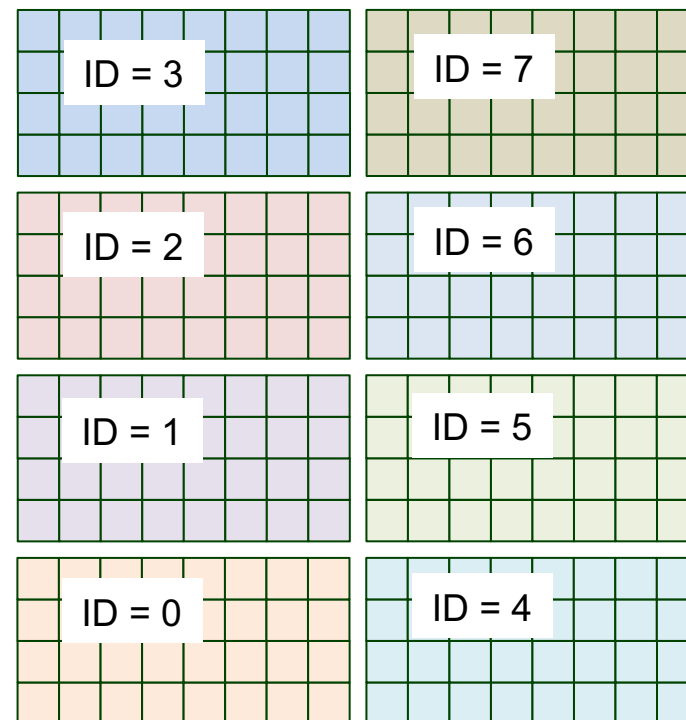
Mapping domains to processes (III)

► Metis partitioning



Graph File:

```
7 11
5 3 2
1 3 4
5 4 2 1
2 3 6 7
1 3 6
5 4 7
6 4
```



```
8 10
1 4
2 5 0
3 6 1
7 2
0 5
4 1 6
5 2 7
6 3
```



Exercise – FDM problem (Stencil)

► 1D : 3-point stencil


$$u_i = f(u_{i-1}, u_i, u_{i+1})$$

► 2D : 5-point stencil

$$u_{i,j} = f(u_{i-1,j}, u_{i,j}, u_{i+1,j}, u_{i,j-1}, u_{i,j+1})$$

► 3D : 7-point stencil

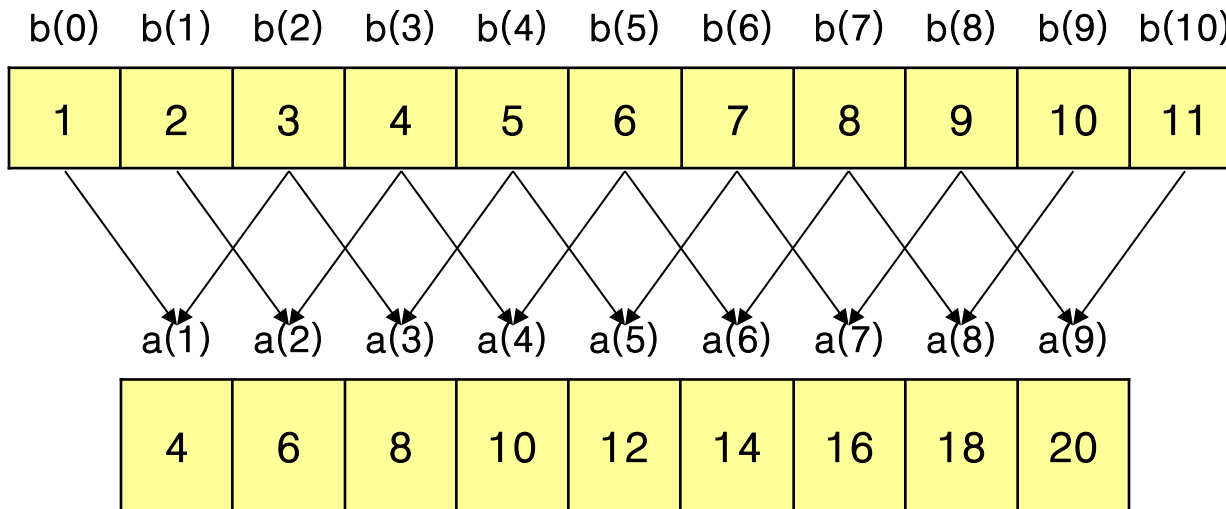
$$u_{i,j,k} = f(u_{i-1,j,k}, u_{i,j,k}, u_{i+1,j,k}, u_{i,j-1,k}, u_{i,j+1,k}, u_{i,j,k-1}, u_{i,j,k+1})$$





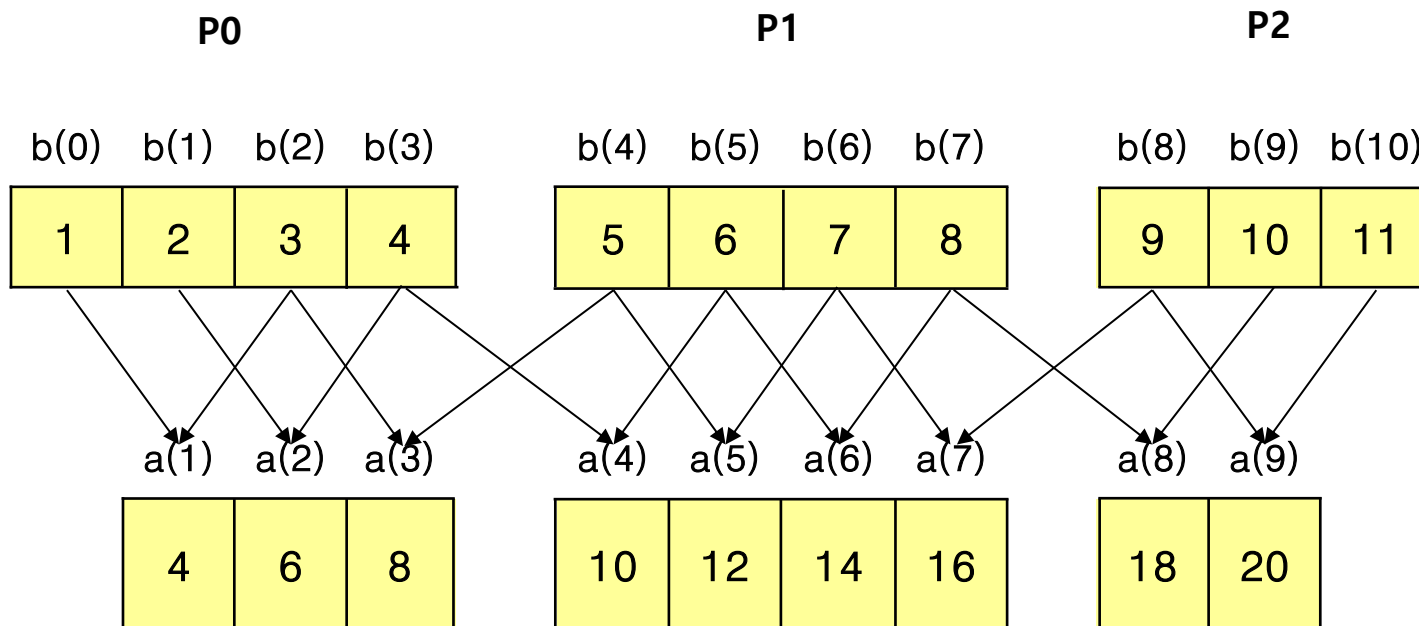
1D FDM: Compute pattern

```
for i in range (2, 11)
    a(i) = b(i-1) + b(i+1)
```






1D FDM: Domain decomposition







1D FDM: Data Dependence and Movements

Process 0



send

receive



ista2 iend1

Process 1




send

send

receive

receive





ista2 iend1

Process 2



send

receive



ista2 iend1



Six-steps of domain decomposition

► Parallelization steps

1. Break up the domain into blocks (domain).
2. Assign blocks to MPI processes one-to-one.
3. Provide a "map" of neighbors to each process.

MPI setup

4. Insert communication subroutine calls where needed.

Communication

5. Write or modify your code so it only updates a single block.
6. Adjust the boundary conditions code.



Serial code

```
import numpy as np

n = 11

a = np.zeros(n, dtype = np.int32)
b = np.zeros(n, dtype = np.int32)

for i in range(0, n) :
    b[i] = i + 1

for i in range(1, n-1) :
    a[i] = b[i-1] + b[i+1]

print(b)
print(a)
```



Step1. Break up the domain

b(0)	b(1)	b(2)	b(3)
1	2	3	4

b(4)	b(5)	b(6)	b(7)
5	6	7	8

b(8)	b(9)	b(10)
9	10	11

a(1)	a(2)	a(3)
4	6	8

a(4)	a(5)	a(6)	a(7)
10	12	14	16

a(8)	a(9)
18	20



Step2. Assign blocks to MPI processes one-to-one

P0

b(0)	b(1)	b(2)	b(3)
1	2	3	4

a(1)	a(2)	a(3)
4	6	8

ista_b = 1, iend_b = 4
ista_a = 2, iend_a = 4

P1

b(4)	b(5)	b(6)	b(7)
5	6	7	8

a(4)	a(5)	a(6)	a(7)
10	12	14	16

ista_b = 5, iend_b = 8
ista_a = 5, iend_a = 8

P2

b(8)	b(9)	b(10)
9	10	11

a(8)	a(9)
18	20

ista_b = 9, iend_b = 11
ista_a = 9, iend_a = 10

```
ista_b, iend_b = para_range(0, n - 1, size, rank)
ista_a = ista_b; iend_a = iend_b

if rank == 0 :
    ista_a = 1
if rank == size - 1 :
    iend_a = n - 2
```



Step3. Provide a "map" of neighbors to each process

P0

b(0) b(1) b(2) b(3)

1	2	3	4
---	---	---	---

a(1) a(2) a(3)

4	6	8
---	---	---

p_next = rank+1
p_prev = MPI_PROC_NULL

P1

b(4) b(5) b(6) b(7)

5	6	7	8
---	---	---	---

a(4) a(5) a(6) a(7)

10	12	14	16
----	----	----	----

p_next = rank+1
p_prev = rank-1

P2

b(8) b(9) b(10)

9	10	11
---	----	----

a(8) a(9)

18	20
----	----

p_next = MPI_PROC_NULL
p_prev = rank-1

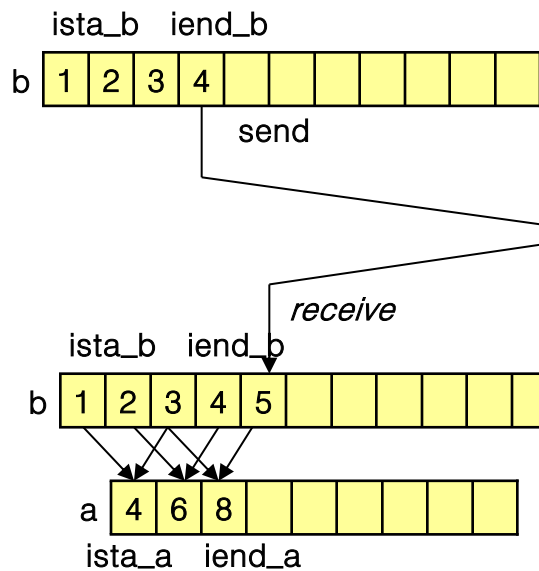
```
p_next = rank + 1; p_prev = rank - 1
```

```
if rank == size - 1 :  
    p_next = MPI.PROC_NULL  
if rank == 0 :  
    p_prev = MPI.PROC_NULL
```





Step 4. Insert communication subroutine calls

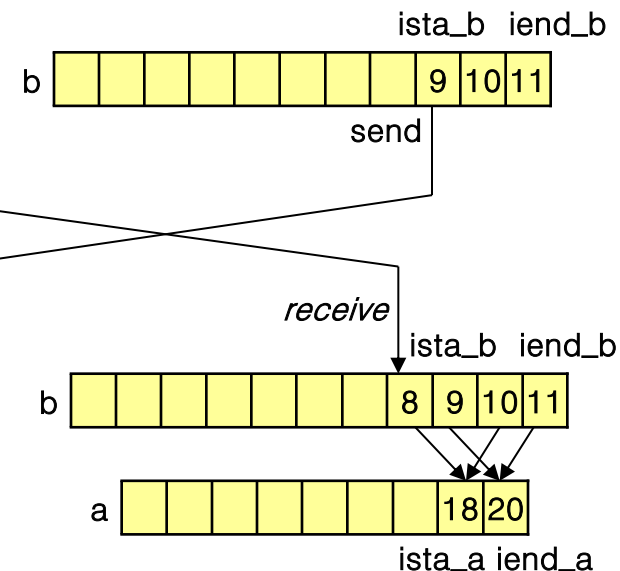
Process 0



Process 1



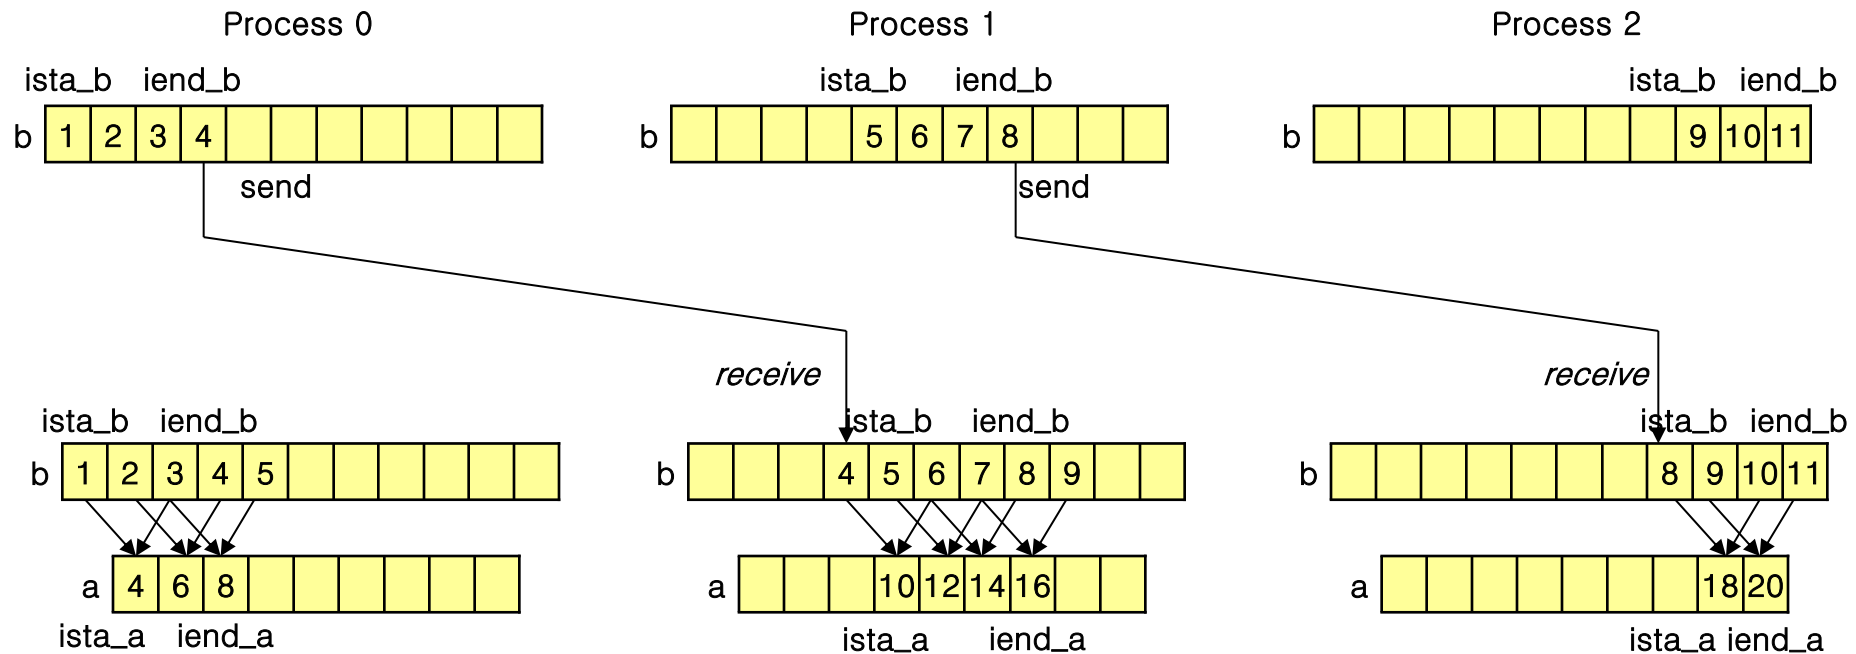
Process 2





Step 4. Insert communication subroutine calls

Stage 1. send to p_next, receive from p_prev



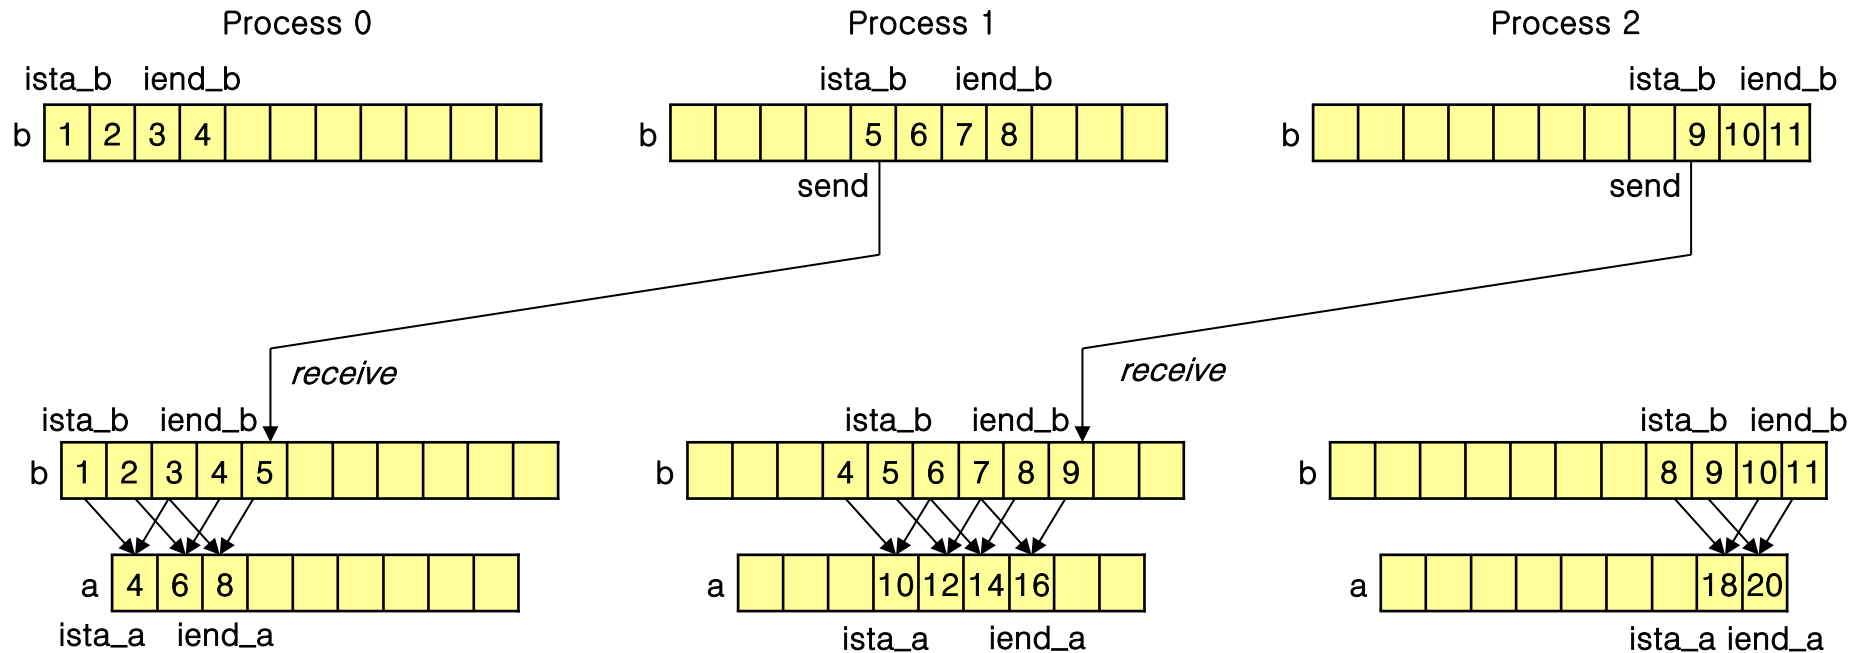
```
req_i1 = comm.Isend(b[iend:iend+1], inext, 11)
req_r1 = comm.Irecv(b[ista-1: ista], iprev, 11)

MPI.Request.Wait(req_i1)
MPI.Request.Wait(req_r1)
```



Step 4. Insert communication subroutine calls

Stage 2. send to p_prev, receive from p_next




```
req_i2 = comm.Isend(b[ista:ista+1], iprev, 12)
req_r2 = comm.Irecv(b[iend+1: iend+2], inext, 12)

MPI.Request.Wait(req_i2)
MPI.Request.Wait(req_r2)
```

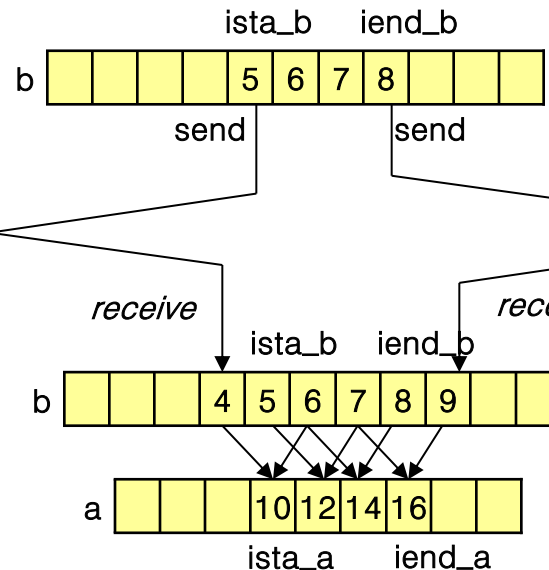


Step 5. Modify a code so it only updates a single block.

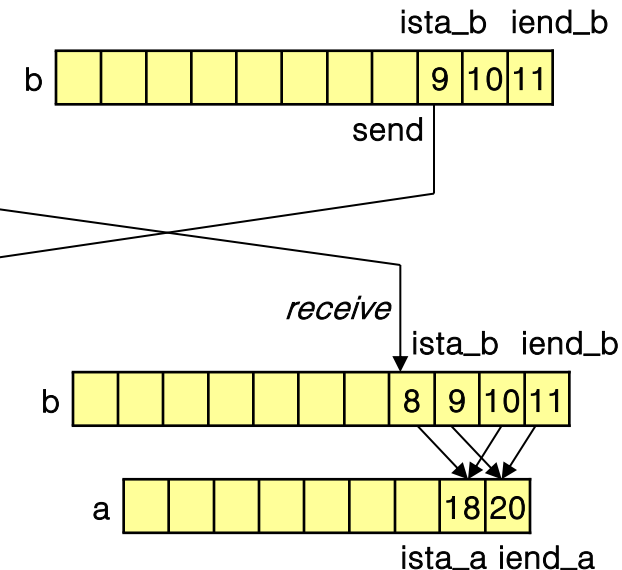
Process 0



Process 1



Process 2



```
for i in range(ista_b, iend_b+1) :  
    b[i] = i + 1  
  
# Communications  
  
for i in range(ista_a, iend_a+1) :  
    a[i] = b[i-1] + b[i+1]
```



Parallel implementation (I)

```
from mpi4py import MPI
import numpy as np
```

```
comm = MPI.COMM_WORLD
```

```
size = comm.Get_size()
rank = comm.Get_rank()
```

```
n = 11
```

```
a = np.zeros(n, dtype = np.int32)
b = np.zeros(n, dtype = np.int32)
```

```
ista_b, iend_b = para_range(0, n - 1, size, rank)
ista_a = ista_b; iend_a = iend_b
```

```
if rank == 0 :
    ista_a = 1
if rank == size - 1 :
    iend_a = n - 2
```

Step2

```
p_next = rank + 1; p_prev = rank - 1
```

```
if rank == size - 1 :
    p_next = MPI.PROC_NULL
if rank == 0 :
    p_prev = MPI.PROC_NULL
```

Step3



Parallel implementation (II)

```
for i in range(ista_b, iend_b+1) :  
    b[i] = i + 1
```

```
req_i1 = comm.Isend(b[iend_b:iend_b+1], p_next, 11)  
req_i2 = comm.Isend(b[ista_b:ista_b+1], p_prev, 12)  
req_r1 = comm.Irecv(b[ista_b-1: ista_b], p_prev, 11)  
req_r2 = comm.Irecv(b[iend_b+1: iend_b+2], p_next, 12)
```

```
MPI.Request.Wait(req_i1)  
MPI.Request.Wait(req_i2)  
MPI.Request.Wait(req_r1)  
MPI.Request.Wait(req_r2)
```

Step4

```
for i in range(ista_a, iend_a+1) :  
    a[i] = b[i-1] + b[i+1]
```

```
for i in range(size) :  
    if i == rank :  
        print(rank)  
        print(b)  
        print(a)  
comm.Barrier()
```

Step5

Where is Step6?

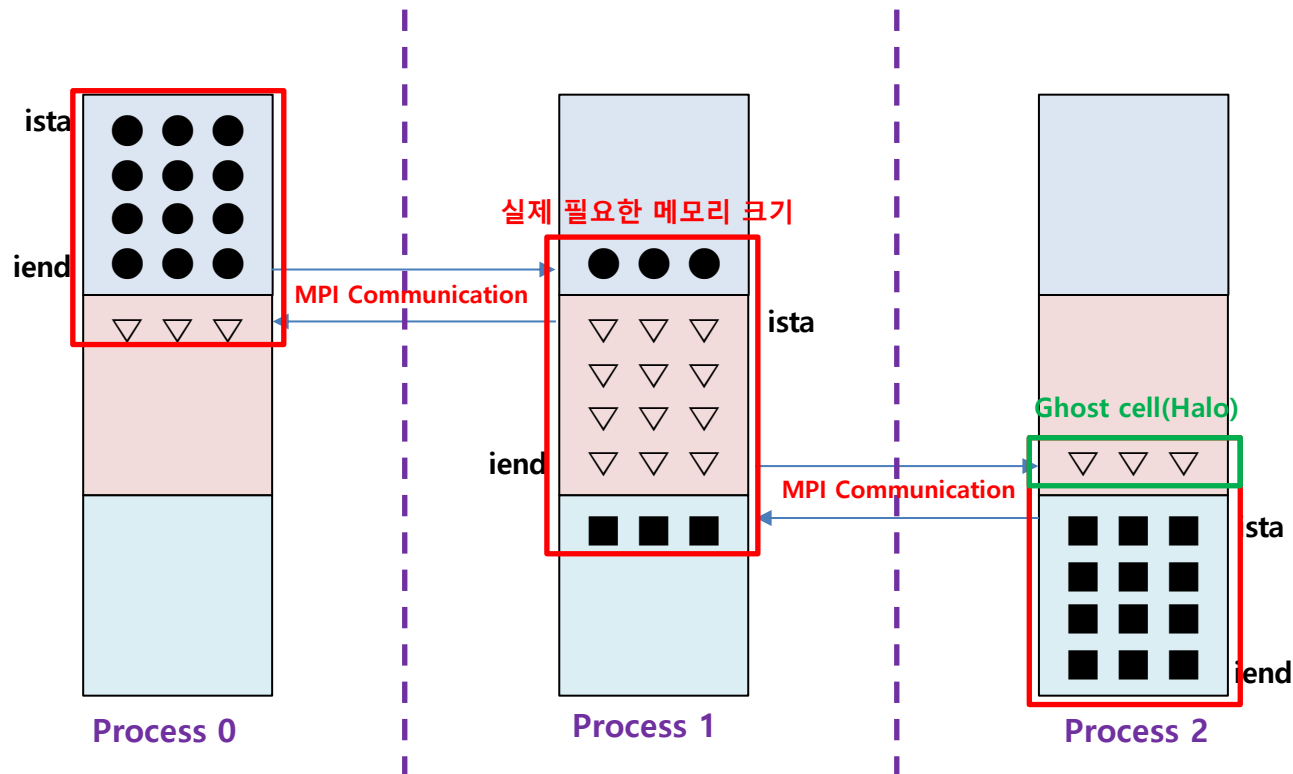
It is incorporated in ista_a and iend_a



2D FDM : Row-Wise

► Row-Wise Decomposition

- C/Python language





2D FDM : Column-Wise

► Column-wise Decomposition

- Use temporary buffer because data in memory are not continuous.

